

```
from google.colab import files
import pandas as pd
uploaded = files.upload()
df = pd.read_csv("dataset.csv")
```

<IPython.core.display.HTML object>

Saving dataset.csv to dataset.csv

```
print(df.head())
```

	Date	Location	Latitude	Longitude	SST	pH
Level \						
0	01-01-2015	Red Sea	20.0248	38.4931	29.47	8.107
1	07-01-2015	Great Barrier Reef	-18.2988	147.7782	29.65	8.004
2	14-01-2015	Caribbean Sea	14.9768	-75.0233	28.86	7.947
3	20-01-2015	Great Barrier Reef	-18.3152	147.6486	28.97	7.995
4	27-01-2015	Galápagos	-0.8805	-90.9769	28.60	7.977

	Bleaching Severity	Species Observed	Marine Heatwave
0	NaN	106	False
1	High	116	False
2	High	90	False
3	Medium	94	False
4	NaN	110	False

Implementing Logistic Regression

```
# STEP 0: Imports
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split

# STEP 1: Load the dataset
from google.colab import files
uploaded = files.upload()

# Replace 'dataset.csv' with your filename if needed
df = pd.read_csv("dataset.csv")

# STEP 2: Create binary target
```

```
df['Collapse'] = df['Species Observed'].apply(lambda x: 1 if x < 100
else 0)
```

```
# STEP 3: Drop unnecessary columns
```

```
df.drop(['Date', 'Species Observed'], axis=1, inplace=True)
```

```
# STEP 4: Encode categorical features
```

```
le_location = LabelEncoder()
```

```
df['Location'] = le_location.fit_transform(df['Location'])
```

```
le_bleaching = LabelEncoder()
```

```
df['Bleaching Severity'] = le_bleaching.fit_transform(df['Bleaching
Severity'])
```

```
df['Marine Heatwave'] = df['Marine Heatwave'].astype(int)
```

```
# STEP 5: Train-test split
```

```
X = df.drop('Collapse', axis=1)
```

```
y = df['Collapse']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

```
# STEP 6: Logistic Regression
```

```
clf = LogisticRegression(max_iter=1000)
```

```
clf.fit(X_train, y_train)
```

```
# STEP 7: Feature importance
```

```
importance_lr = pd.DataFrame({
```

```
    'Feature': X.columns,
```

```
    'Importance': clf.coef_[0]
```

```
}).sort_values(by='Importance', key=abs, ascending=False)
```

```
print("\nFeature Importance Table:\n")
```

```
print(importance_lr)
```

```
# STEP 8: Plotting
```

```
plt.figure(figsize=(10, 6))
```

```
sns.barplot(x='Importance', y='Feature', data=importance_lr,
palette='coolwarm')
```

```
plt.title('Feature Importance from Logistic Regression')
```

```
plt.xlabel('Coefficient Value (Importance)')
```

```
plt.ylabel('Feature')
```

```
plt.axvline(0, color='black', linewidth=0.8)
```

```
plt.tight_layout()
```

```
plt.show()
```

<IPython.core.display.HTML object>

Saving dataset.csv to dataset (1).csv

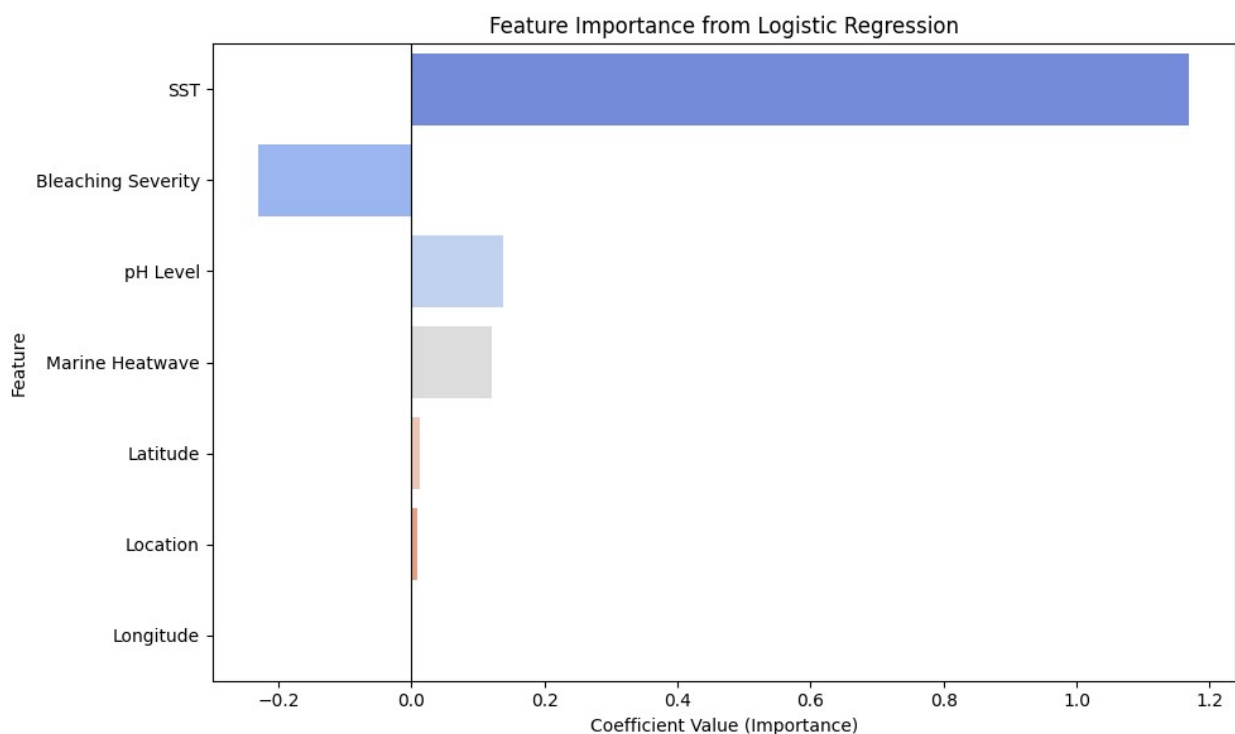
Feature Importance Table:

	Feature	Importance
3	SST	1.169082
5	Bleaching Severity	-0.229563
4	pH Level	0.137481
6	Marine Heatwave	0.121630
1	Latitude	0.013531
0	Location	0.009618
2	Longitude	0.000165

```
/tmp/ipython-input-4-1361258406.py:52: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `y` variable to `hue` and set
`legend=False` for the same effect.
```

```
sns.barplot(x='Importance', y='Feature', data=importance_lr,
palette='coolwarm')
```



Implementing Random Forest Classifier

```
# STEP 0: Imports
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.ensemble import RandomForestClassifier
```

```

from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split

# STEP 1: Load the dataset
from google.colab import files
uploaded = files.upload()

# Replace 'dataset.csv' with your filename if needed
df = pd.read_csv("dataset.csv")

# STEP 2: Create binary target
df['Collapse'] = df['Species Observed'].apply(lambda x: 1 if x < 100
else 0)

# STEP 3: Drop unnecessary columns
df.drop(['Date', 'Species Observed'], axis=1, inplace=True)

# STEP 4: Encode categorical features
le_location = LabelEncoder()
df['Location'] = le_location.fit_transform(df['Location'])

le_bleaching = LabelEncoder()
df['Bleaching Severity'] = le_bleaching.fit_transform(df['Bleaching
Severity'])

df['Marine Heatwave'] = df['Marine Heatwave'].astype(int)

# STEP 5: Train-test split
X = df.drop('Collapse', axis=1)
y = df['Collapse']
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# STEP 6: Random Forest Classifier
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)

# STEP 7: Feature importance
importances = rf.feature_importances_
importance_rf = pd.DataFrame({
    'Feature': X.columns,
    'Importance': importances
}).sort_values(by='Importance', ascending=False)

print("\nFeature Importance Table:\n")
print(importance_rf)

# STEP 8: Plotting
plt.figure(figsize=(10, 6))
sns.barplot(x='Importance', y='Feature', data=importance_rf,

```

```
palette='viridis')
plt.title('Feature Importance from Random Forest Classifier')
plt.xlabel('Feature Importance (Gini Importance)')
plt.ylabel('Feature')
plt.tight_layout()
plt.show()
```

<IPython.core.display.HTML object>

Saving dataset.csv to dataset (2).csv

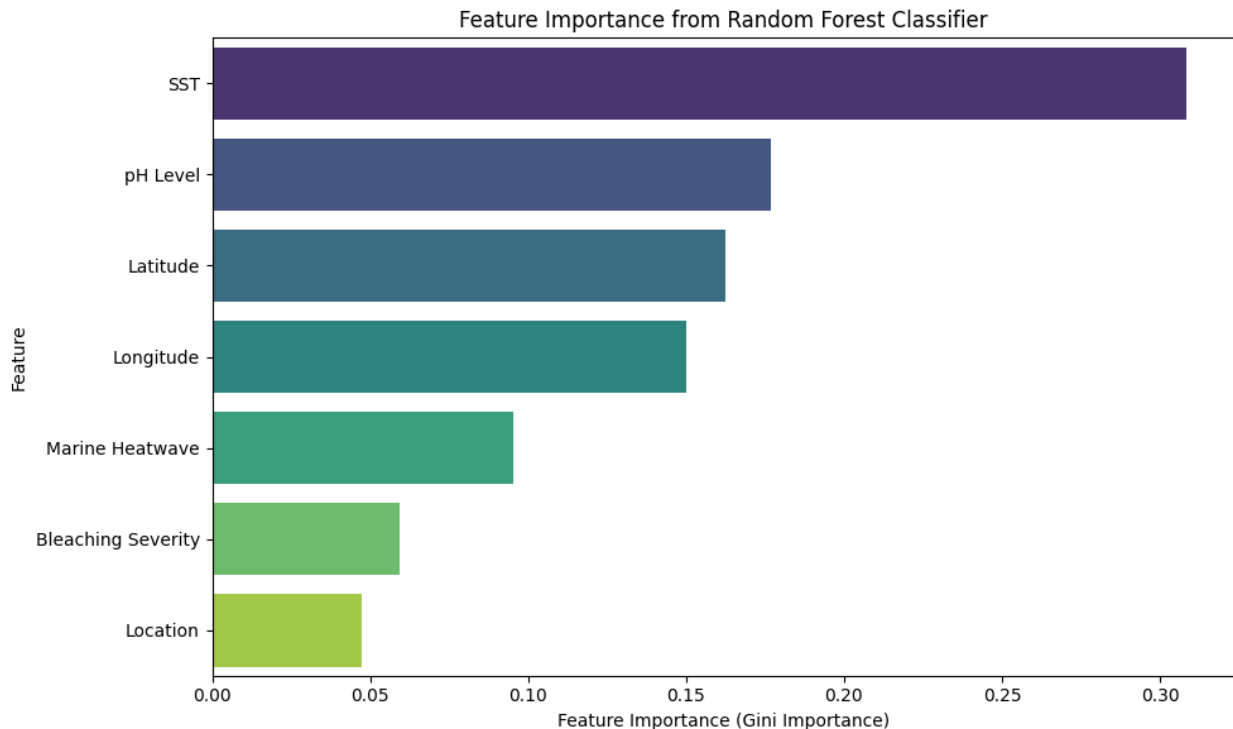
Feature Importance Table:

	Feature	Importance
3	SST	0.308541
4	pH Level	0.177047
1	Latitude	0.162516
2	Longitude	0.150043
6	Marine Heatwave	0.095202
5	Bleaching Severity	0.059225
0	Location	0.047427

/tmp/ipython-input-5-2325563287.py:53: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='Importance', y='Feature', data=importance_rf,
palette='viridis')
```



Implementing Linear SVM

```
# STEP 0: Imports
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.svm import LinearSVC
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

# STEP 1: Load the dataset
from google.colab import files
uploaded = files.upload()

# Replace 'dataset.csv' with your filename if needed
df = pd.read_csv("dataset.csv")

# STEP 2: Create binary target
df['Collapse'] = df['Species Observed'].apply(lambda x: 1 if x < 100
else 0)

# STEP 3: Drop unnecessary columns
df.drop(['Date', 'Species Observed'], axis=1, inplace=True)

# STEP 4: Encode categorical features
le_location = LabelEncoder()
```

```

df['Location'] = le_location.fit_transform(df['Location'])

le_bleaching = LabelEncoder()
df['Bleaching Severity'] = le_bleaching.fit_transform(df['Bleaching
Severity'])

df['Marine Heatwave'] = df['Marine Heatwave'].astype(int)

# STEP 5: Train-test split
X = df.drop('Collapse', axis=1)
y = df['Collapse']
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# STEP 6: Standardization (important for SVM)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# STEP 7: Linear SVM
svm_clf = LinearSVC(max_iter=10000, random_state=42)
svm_clf.fit(X_train_scaled, y_train)

# STEP 8: Feature importance via coefficients
importance_svm = pd.DataFrame({
    'Feature': X.columns,
    'Importance': svm_clf.coef_[0]
}).sort_values(by='Importance', key=abs, ascending=False)

print("\nFeature Importance Table (Linear SVM):\n")
print(importance_svm)

# STEP 9: Plotting
plt.figure(figsize=(10, 6))
sns.barplot(x='Importance', y='Feature', data=importance_svm,
palette='rocket')
plt.title('Feature Importance from Linear SVM')
plt.xlabel('Coefficient Value (Importance)')
plt.ylabel('Feature')
plt.axvline(0, color='black', linestyle='--')
plt.tight_layout()
plt.show()

```

<IPython.core.display.HTML object>

Saving dataset.csv to dataset (3).csv

Feature Importance Table (Linear SVM):

	Feature	Importance
3	SST	0.348321

```

6 Marine Heatwave 0.118440
5 Bleaching Severity -0.065649
1 Latitude 0.041196
4 pH Level 0.031176
2 Longitude 0.009017
0 Location 0.002789

```

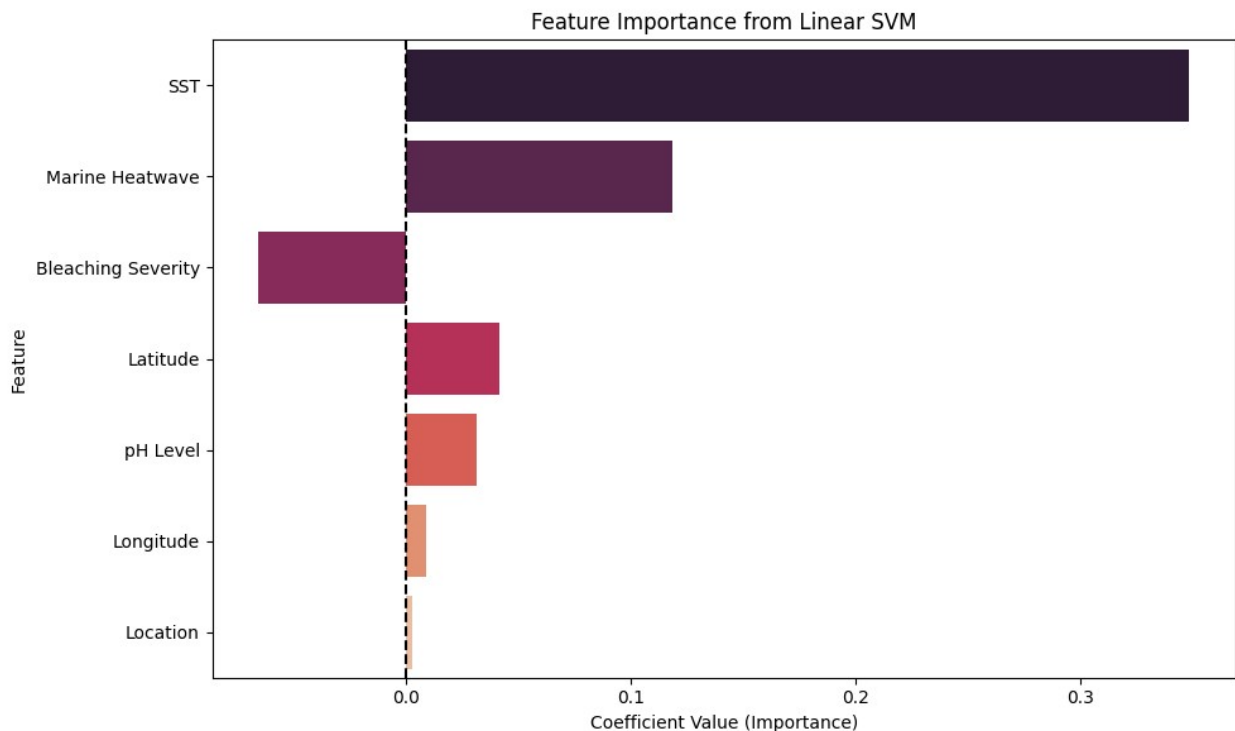
/tmp/ipython-input-6-2309594835.py:58: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```

sns.barplot(x='Importance', y='Feature', data=importance_svm,
palette='rocket')

```



Implementing XGBoost Classifier

```

# STEP 1: Imports
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from xgboost import XGBClassifier, plot_importance
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split

```



```

# STEP 2: Load dataset
from google.colab import files
uploaded = files.upload()

# Replace with your actual filename if needed
df = pd.read_csv("dataset.csv")

# STEP 3: Create binary target
df['Collapse'] = df['Species Observed'].apply(lambda x: 1 if x < 100
else 0)

# STEP 4: Drop unnecessary columns
df.drop(['Date', 'Species Observed'], axis=1, inplace=True)

# STEP 5: Encode categorical features
le_location = LabelEncoder()
df['Location'] = le_location.fit_transform(df['Location'])

le_bleaching = LabelEncoder()
df['Bleaching Severity'] = le_bleaching.fit_transform(df['Bleaching
Severity'])

df['Marine Heatwave'] = df['Marine Heatwave'].astype(int)

# STEP 6: Train-test split
X = df.drop('Collapse', axis=1)
y = df['Collapse']
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# STEP 7: Train XGBoost Classifier
xgb_clf = XGBClassifier(use_label_encoder=False,
eval_metric='logloss', random_state=42)
xgb_clf.fit(X_train, y_train)

# STEP 8: Feature importance
importance_xgb = pd.DataFrame({
    'Feature': X.columns,
    'Importance': xgb_clf.feature_importances_
}).sort_values(by='Importance', ascending=False)

print("\nFeature Importance Table (XGBoost):\n")
print(importance_xgb)

# STEP 9: Plotting
plt.figure(figsize=(10, 6))
sns.barplot(x='Importance', y='Feature', data=importance_xgb,
palette='crest')
plt.title('Feature Importance from XGBoost Classifier')
plt.xlabel('Feature Importance Score')

```

```
plt.ylabel('Feature')
plt.tight_layout()
plt.show()
```

<IPython.core.display.HTML object>

Saving dataset.csv to dataset (4).csv

Feature Importance Table (XGBoost):

	Feature	Importance
3	SST	0.375568
5	Bleaching Severity	0.135949
1	Latitude	0.127916
4	pH Level	0.127279
2	Longitude	0.122053
0	Location	0.111235
6	Marine Heatwave	0.000000

/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158:

UserWarning: [13:13:24] WARNING: /workspace/src/learner.cc:740:

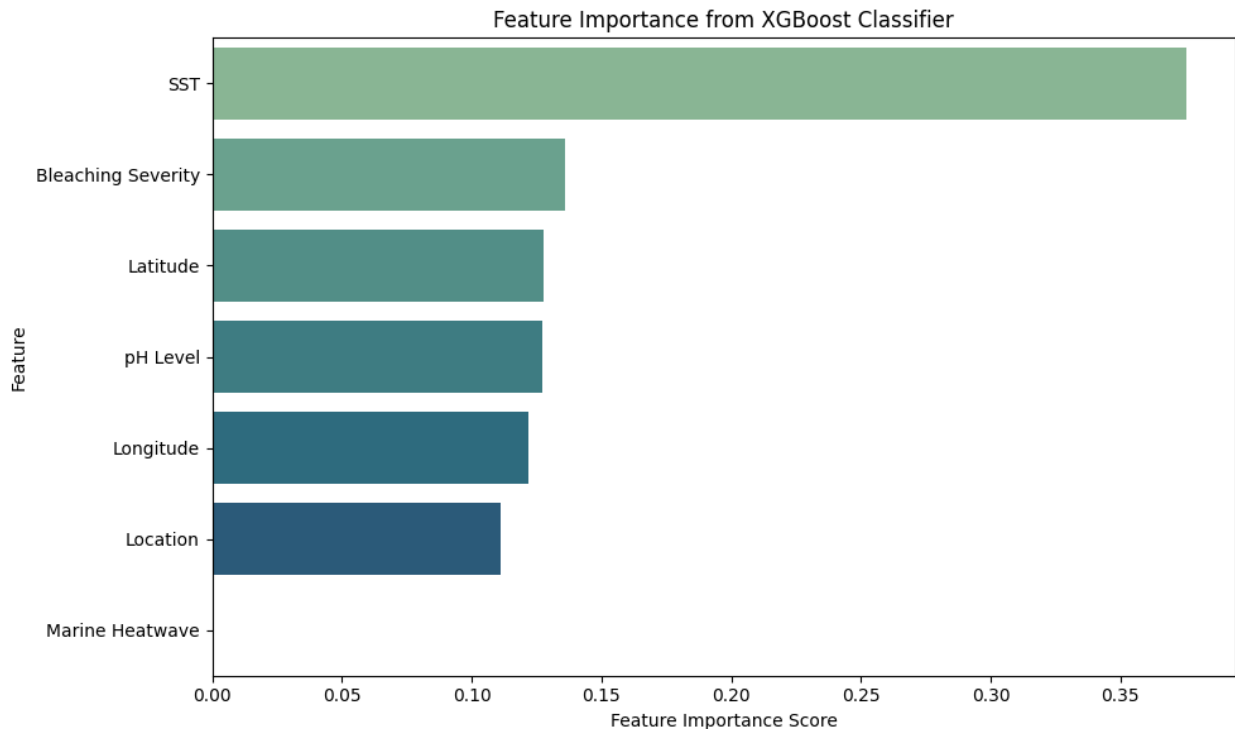
Parameters: { "use_label_encoder" } are not used.

```
warnings.warn(smsg, UserWarning)
```

/tmp/ipython-input-7-2343520320.py:52: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='Importance', y='Feature', data=importance_xgb,
palette='crest')
```



Implementing KNN Classifier

```
# STEP 1: Imports
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.inspection import permutation_importance

# STEP 2: Load dataset
from google.colab import files
uploaded = files.upload()

# Load CSV
df = pd.read_csv("dataset.csv")

# STEP 3: Create binary target
df['Collapse'] = df['Species Observed'].apply(lambda x: 1 if x < 100
else 0)

# STEP 4: Drop unused columns
df.drop(['Date', 'Species Observed'], axis=1, inplace=True)

# STEP 5: Encode categorical features
```

```

le_location = LabelEncoder()
df['Location'] = le_location.fit_transform(df['Location'])

le_bleaching = LabelEncoder()
df['Bleaching Severity'] = le_bleaching.fit_transform(df['Bleaching
Severity'])

df['Marine Heatwave'] = df['Marine Heatwave'].astype(int)

# STEP 6: Train-test split
X = df.drop('Collapse', axis=1)
y = df['Collapse']
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# STEP 7: Scale features (very important for KNN)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# STEP 8: Train KNN Classifier
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train_scaled, y_train)

# STEP 9: Permutation Feature Importance (since KNN doesn't have
coefficients)
perm_importance = permutation_importance(knn, X_test_scaled, y_test,
n_repeats=20, random_state=42)

# STEP 10: Format importance into DataFrame
importance_knn = pd.DataFrame({
    'Feature': X.columns,
    'Importance': perm_importance.importances_mean
}).sort_values(by='Importance', ascending=False)

print("\nFeature Importance Table (KNN via Permutation):\n")
print(importance_knn)

# STEP 11: Plotting
plt.figure(figsize=(10, 6))
sns.barplot(x='Importance', y='Feature', data=importance_knn,
palette='flare')
plt.title('Feature Importance from KNN Classifier (Permutation)')
plt.xlabel('Mean Importance (Decrease in Accuracy)')
plt.ylabel('Feature')
plt.tight_layout()
plt.show()

<IPython.core.display.HTML object>

```

```
Saving dataset.csv to dataset (5).csv
```


```
Feature Importance Table (KNN via Permutation):
```

	Feature	Importance
6	Marine Heatwave	0.1005
3	SST	0.0390
5	Bleaching Severity	0.0250
0	Location	0.0220
2	Longitude	0.0215
1	Latitude	0.0180
4	pH Level	0.0145

```
/tmp/ipython-input-8-210698949.py:61: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.
```

```
sns.barplot(x='Importance', y='Feature', data=importance_knn, palette='flare')
```



Feature Importance Comparison across all models

```
# STEP 0: Prepare individual importance DataFrames (you should already have them)  
# Each must have: ['Feature', 'Importance'] columns
```

```

# e.g. importance_lr, importance_rf, etc. are your earlier outputs

# STEP 1: Combine Importances
all_importances = [importance_lr, importance_rf, importance_svm,
importance_knn, importance_xgb]
model_names = ['Logistic Regression', 'Random Forest', 'Linear SVM',
'KNN', 'XGBoost']

# Merge into one DataFrame
feature_importance_df = pd.DataFrame()

for imp_df, name in zip(all_importances, model_names):
    feature_importance_df[name] = imp_df.set_index('Feature')
['Importance']

# Align indexes and fill missing with 0
feature_importance_df = feature_importance_df.fillna(0)
feature_importance_df['Mean_Importance'] =
feature_importance_df.mean(axis=1)

# STEP 2: Normalize for 0-1 scale comparison
normalized_df =
feature_importance_df.drop(columns=["Mean_Importance"])
normalized_df = normalized_df / normalized_df.max()

# STEP 3: Plotting with Line Graph and Annotations
import matplotlib.pyplot as plt

plt.figure(figsize=(12, 6))

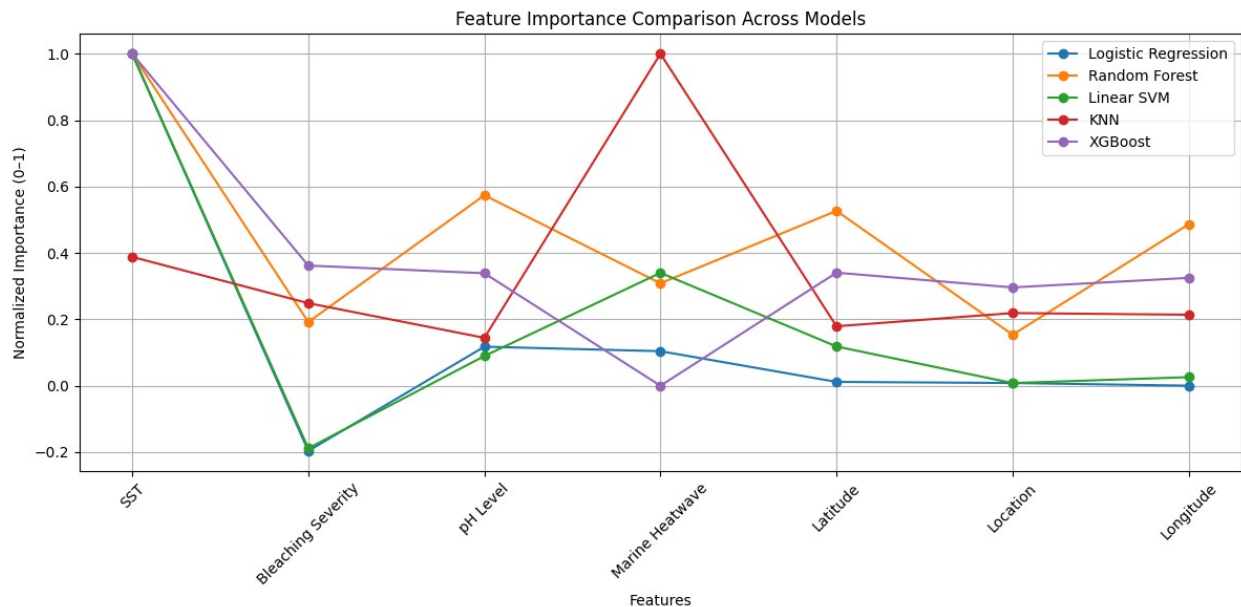
# Plot each model's line and annotate each point
for model in normalized_df.columns:
    plt.plot(normalized_df.index, normalized_df[model], marker='o',
label=model)

    # Annotate each point for this model
    '''for i, value in enumerate(normalized_df[model]):
        plt.annotate(
            f"{value:.2f}",
            (normalized_df.index[i], value),
            textcoords="offset points",
            xytext=(0, 5), # vertical offset
            ha='center',
            fontsize=8,
            color='black'
        )'''

plt.xticks(rotation=45)
plt.xlabel("Features")
plt.ylabel("Normalized Importance (0-1)")

```

```
plt.title("Feature Importance Comparison Across Models")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```



Calculating Evaluation Metrics using XGBoost model and SST, Bleaching Severity, Location and Longitude

```
# STEP 0: Re-import everything (in case of runtime reset)
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score
from sklearn.preprocessing import LabelEncoder
from xgboost import XGBClassifier

# STEP 1: Load the dataset
from google.colab import files
uploaded = files.upload()

df = pd.read_csv("dataset.csv") # Make sure the file name is correct

# STEP 2: Check columns
print("Available columns:", df.columns.tolist())

# STEP 3: Create binary target column
if 'Species Observed' in df.columns:
    df['Collapse'] = df['Species Observed'].apply(lambda x: 1 if x <
100 else 0)
```

```

else:
    raise KeyError("[] Column 'Species Observed' not found in the
uploaded dataset.")

# STEP 4: Encode categorical features
le_location = LabelEncoder()
df['Location'] = le_location.fit_transform(df['Location'])

le_bleaching = LabelEncoder()
df['Bleaching Severity'] = le_bleaching.fit_transform(df['Bleaching
Severity'])

df['Marine Heatwave'] = df['Marine Heatwave'].astype(int)

# STEP 5: Select required features
features = ['SST', 'Bleaching Severity', 'Location', 'Longitude']
X = df[features]
y = df['Collapse']

# STEP 6: Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# STEP 7: Train and Evaluate XGBoost Classifier
model = XGBClassifier(use_label_encoder=False, eval_metric='logloss',
random_state=42)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)

# STEP 8: Metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print("\n[] Evaluation Metrics:")
print(f"Accuracy : {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall    : {recall:.2f}")
print(f"F1 Score  : {f1:.2f}")

<IPython.core.display.HTML object>

Saving dataset.csv to dataset (13).csv
Available columns: ['Date', 'Location', 'Latitude', 'Longitude',
'SST', 'pH Level', 'Bleaching Severity', 'Species Observed', 'Marine
Heatwave']

/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158:
UserWarning: [06:53:27] WARNING: /workspace/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

```



```
warnings.warn(smsg, UserWarning)
```

□ Evaluation Metrics:

Accuracy : 0.84

Precision: 0.70

Recall : 0.35

F1 Score : 0.47

Calculating Evaluation Metrics using Random Forest model and SST, Bleaching Severity, Location and Longitude

```
from sklearn.ensemble import RandomForestClassifier

from google.colab import files
uploaded = files.upload()

# Recreate the binary target
df['Collapse'] = df['Species Observed'].apply(lambda x: 1 if x < 100
else 0)

# Encode categorical variables
le_location = LabelEncoder()
df['Location'] = le_location.fit_transform(df['Location'])

le_bleaching = LabelEncoder()
df['Bleaching Severity'] = le_bleaching.fit_transform(df['Bleaching
Severity'])

df['Marine Heatwave'] = df['Marine Heatwave'].astype(int)

# Select the same parameters as before
features = ['SST', 'Bleaching Severity', 'Location', 'Longitude']
X = df[features]
y = df['Collapse']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Train Random Forest Classifier
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)
y_pred_rf = rf_model.predict(X_test)

# Calculate metrics
accuracy_rf = accuracy_score(y_test, y_pred_rf)
precision_rf = precision_score(y_test, y_pred_rf)
recall_rf = recall_score(y_test, y_pred_rf)
```

```
f1_rf = f1_score(y_test, y_pred_rf)

print("\n Evaluation Metrics:")
print(f"Accuracy : {accuracy_rf:.2f}")
print(f"Precision: {precision_rf:.2f}")
print(f"Recall    : {recall_rf:.2f}")
print(f"F1 Score  : {f1_rf:.2f}")
```

<IPython.core.display.HTML object>

Saving dataset.csv to dataset (15).csv

```
 Evaluation Metrics:
Accuracy : 0.86
Precision: 0.80
Recall    : 0.40
F1 Score  : 0.53
```

Calculating Evaluation Metrics using Random Forest model and SST

```
# STEP 0: Imports
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score

# STEP 1: Load dataset
from google.colab import files
uploaded = files.upload()

df = pd.read_csv("dataset.csv") # Replace if filename differs

# STEP 2: Create binary target
df['Collapse'] = df['Species Observed'].apply(lambda x: 1 if x < 100
else 0)

# STEP 3: Select SST as the only feature
X = df[['SST']]
y = df['Collapse']

# STEP 4: Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# STEP 5: Train Random Forest Classifier
model = RandomForestClassifier(n_estimators=100, random_state=42)
model.fit(X_train, y_train)

# STEP 6: Predict and Evaluate
```

```

y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print("\n Evaluation Metrics using only SST:")
print(f"Accuracy : {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall    : {recall:.2f}")
print(f"F1 Score  : {f1:.2f}")

```

<IPython.core.display.HTML object>

Saving dataset.csv to dataset (16).csv

```

Evaluation Metrics using only SST:
Accuracy : 0.80
Precision: 0.50
Recall    : 0.40
F1 Score  : 0.44

```

Calculating Evaluation Metrics using Random Forest model and SST and Bleaching Severity

```

# STEP 0: Imports
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score

# STEP 1: Upload dataset
from google.colab import files
uploaded = files.upload()

# STEP 2: Load and preprocess
df = pd.read_csv("dataset.csv")

# Create binary target
df['Collapse'] = df['Species Observed'].apply(lambda x: 1 if x < 100
else 0)

# Encode Bleaching Severity
le_bleaching = LabelEncoder()
df['Bleaching Severity'] = le_bleaching.fit_transform(df['Bleaching
Severity'])

# STEP 3: Select features and target

```

```

X = df[['SST', 'Bleaching Severity']]
y = df['Collapse']

# STEP 4: Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# STEP 5: Random Forest Classifier
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)

# STEP 6: Predictions
y_pred = rf.predict(X_test)

# STEP 7: Evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print("\n Evaluation using SST + Bleaching Severity:")
print(f"Accuracy : {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall    : {recall:.2f}")
print(f"F1 Score  : {f1:.2f}")

<IPython.core.display.HTML object>

Saving dataset.csv to dataset (17).csv

Evaluation using SST + Bleaching Severity:
Accuracy : 0.87
Precision: 0.73
Recall    : 0.55
F1 Score  : 0.63

```

Calculating Evaluation Metrics using XGBoost model and SST and Bleaching Severity

```

# STEP 1: Imports
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score
from xgboost import XGBClassifier

# STEP 2: Upload dataset
from google.colab import files
uploaded = files.upload()

```

```

# STEP 3: Load and preprocess
df = pd.read_csv("dataset.csv")

# Create binary target
df['Collapse'] = df['Species Observed'].apply(lambda x: 1 if x < 100
else 0)

# Encode Bleaching Severity
le_bleaching = LabelEncoder()
df['Bleaching Severity'] = le_bleaching.fit_transform(df['Bleaching
Severity'])

# STEP 4: Select features and target
X = df[['SST', 'Bleaching Severity']]
y = df['Collapse']

# STEP 5: Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# STEP 6: XGBoost Classifier
xgb = XGBClassifier(use_label_encoder=False, eval_metric='logloss',
random_state=42)
xgb.fit(X_train, y_train)

# STEP 7: Predictions
y_pred = xgb.predict(X_test)

# STEP 8: Evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print("\n Evaluation using SST + Bleaching Severity:")
print(f"Accuracy : {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall    : {recall:.2f}")
print(f"F1 Score  : {f1:.2f}")

```

<IPython.core.display.HTML object>

Saving dataset.csv to dataset (18).csv

Evaluation using SST + Bleaching Severity:

```

Accuracy : 0.84
Precision: 0.67
Recall    : 0.40
F1 Score  : 0.50

```

```
/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158:
UserWarning: [08:55:36] WARNING: /workspace/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.
```

```
warnings.warn(smsg, UserWarning)
```

Calculating Evaluation Metrics using XGBoost model and SST and Location

```
# STEP 1: Imports
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score
from xgboost import XGBClassifier

# STEP 2: Upload dataset
from google.colab import files
uploaded = files.upload()

# STEP 3: Load and preprocess
df = pd.read_csv("dataset.csv")

# Create binary target
df['Collapse'] = df['Species Observed'].apply(lambda x: 1 if x < 100
else 0)

# Encode Location (categorical)
le_location = LabelEncoder()
df['Location'] = le_location.fit_transform(df['Location'])

# STEP 4: Select features and target
X = df[['SST', 'Location']]
y = df['Collapse']

# STEP 5: Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# STEP 6: XGBoost Classifier
xgb = XGBClassifier(use_label_encoder=False, eval_metric='logloss',
random_state=42)
xgb.fit(X_train, y_train)

# STEP 7: Predictions
y_pred = xgb.predict(X_test)

# STEP 8: Evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
```

```

recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print("\n Evaluation using SST + Location:")
print(f"Accuracy : {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall    : {recall:.2f}")
print(f"F1 Score  : {f1:.2f}")

<IPython.core.display.HTML object>

Saving dataset.csv to dataset (19).csv

 Evaluation using SST + Location:
Accuracy : 0.82
Precision: 0.62
Recall    : 0.25
F1 Score  : 0.36

/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158:
UserWarning: [08:57:07] WARNING: /workspace/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(msg, UserWarning)

```

Calculating Evaluation Metrics using Random Forest model and SST and Location

```

# STEP 0: Imports
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score

# STEP 1: Upload dataset
from google.colab import files
uploaded = files.upload()

# STEP 2: Load and preprocess
df = pd.read_csv("dataset.csv")

# Create binary target
df['Collapse'] = df['Species Observed'].apply(lambda x: 1 if x < 100
else 0)

# Encode Location (categorical)
le_location = LabelEncoder()
df['Location'] = le_location.fit_transform(df['Location'])

```

```

# STEP 3: Select features and target
X = df[['SST', 'Location']]
y = df['Collapse']

# STEP 4: Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# STEP 5: Random Forest Classifier
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)

# STEP 6: Predictions
y_pred = rf.predict(X_test)

# STEP 7: Evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print("\n Evaluation using SST + Location:")
print(f"Accuracy : {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall    : {recall:.2f}")
print(f"F1 Score  : {f1:.2f}")

<IPython.core.display.HTML object>

Saving dataset.csv to dataset (20).csv

 Evaluation using SST + Location:
Accuracy : 0.82
Precision: 0.62
Recall    : 0.25
F1 Score  : 0.36

```

XGBoost(SST+Longitude)

```

# STEP 1: Imports
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score
from xgboost import XGBClassifier

# STEP 2: Upload dataset
from google.colab import files
uploaded = files.upload()

```



```

# STEP 3: Load and preprocess
df = pd.read_csv("dataset.csv")

# Create binary target
df['Collapse'] = df['Species Observed'].apply(lambda x: 1 if x < 100
else 0)

# STEP 4: Select features and target
X = df[['SST', 'Longitude']]
y = df['Collapse']

# STEP 5: Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# STEP 6: XGBoost Classifier
xgb = XGBClassifier(use_label_encoder=False, eval_metric='logloss',
random_state=42)
xgb.fit(X_train, y_train)

# STEP 7: Predictions
y_pred = xgb.predict(X_test)

# STEP 8: Evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print("\n Evaluation using SST + Longitude:")
print(f"Accuracy : {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall    : {recall:.2f}")
print(f"F1 Score  : {f1:.2f}")

<IPython.core.display.HTML object>

Saving dataset.csv to dataset (21).csv

 Evaluation using SST + Longitude:
Accuracy : 0.84
Precision: 0.70
Recall    : 0.35
F1 Score  : 0.47

/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158:
UserWarning: [09:00:33] WARNING: /workspace/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(msg, UserWarning)

```

Random Forest(SST+Longitude)

```
# STEP 0: Imports
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score

# STEP 1: Upload dataset
from google.colab import files
uploaded = files.upload()

# STEP 2: Load and preprocess
df = pd.read_csv("dataset.csv")

# Create binary target
df['Collapse'] = df['Species Observed'].apply(lambda x: 1 if x < 100
else 0)

# STEP 3: Select features and target
X = df[['SST', 'Longitude']]
y = df['Collapse']

# STEP 4: Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# STEP 5: Random Forest Classifier
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)

# STEP 6: Predictions
y_pred = rf.predict(X_test)

# STEP 7: Evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print("\n Evaluation using SST + Longitude:")
print(f"Accuracy : {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall    : {recall:.2f}")
print(f"F1 Score  : {f1:.2f}")

<IPython.core.display.HTML object>

Saving dataset.csv to dataset (22).csv
```

□ Evaluation using SST + Longitude:

Accuracy : 0.83

Precision: 0.71

Recall : 0.25

F1 Score : 0.37

XGBoost(Bleaching + Location)

```
# STEP 1: Imports
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score
from xgboost import XGBClassifier

# STEP 2: Upload dataset
from google.colab import files
uploaded = files.upload()

# STEP 3: Load and preprocess
df = pd.read_csv("dataset.csv")

# Create binary target
df['Collapse'] = df['Species Observed'].apply(lambda x: 1 if x < 100
else 0)

# Encode categorical features
le_bleaching = LabelEncoder()
df['Bleaching Severity'] = le_bleaching.fit_transform(df['Bleaching
Severity'])

le_location = LabelEncoder()
df['Location'] = le_location.fit_transform(df['Location'])

# STEP 4: Select features and target
X = df[['Bleaching Severity', 'Location']]
y = df['Collapse']

# STEP 5: Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# STEP 6: XGBoost Classifier
xgb = XGBClassifier(use_label_encoder=False, eval_metric='logloss',
random_state=42)
xgb.fit(X_train, y_train)

# STEP 7: Predictions
y_pred = xgb.predict(X_test)
```

```
# STEP 8: Evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print("\n Evaluation using Bleaching Severity + Location:")
print(f"Accuracy : {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall    : {recall:.2f}")
print(f"F1 Score  : {f1:.2f}")
```

<IPython.core.display.HTML object>

Saving dataset.csv to dataset (23).csv

Evaluation using Bleaching Severity + Location:

Accuracy : 0.80
Precision: 0.00
Recall : 0.00
F1 Score : 0.00

/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158:
UserWarning: [09:04:33] WARNING: /workspace/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

```
warnings.warn(smsg, UserWarning)
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
```

Random Forest(Bleaching + Location)

```
# STEP 0: Imports
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

# STEP 1: Upload dataset
from google.colab import files
uploaded = files.upload()

# STEP 2: Load and preprocess
```

```

df = pd.read_csv("dataset.csv")

# Create binary target
df['Collapse'] = df['Species Observed'].apply(lambda x: 1 if x < 100
else 0)

# Encode categorical features
le_bleaching = LabelEncoder()
df['Bleaching Severity'] = le_bleaching.fit_transform(df['Bleaching
Severity'])

le_location = LabelEncoder()
df['Location'] = le_location.fit_transform(df['Location'])

# STEP 3: Select features and target
X = df[['Bleaching Severity', 'Location']]
y = df['Collapse']

# STEP 4: Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# STEP 5: Random Forest Classifier
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)

# STEP 6: Predictions
y_pred = rf.predict(X_test)

# STEP 7: Evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print("\n Evaluation using Bleaching Severity + Location:")
print(f"Accuracy : {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall    : {recall:.2f}")
print(f"F1 Score  : {f1:.2f}")

```

<IPython.core.display.HTML object>

Saving dataset.csv to dataset (24).csv

```

 Evaluation using Bleaching Severity + Location:
Accuracy : 0.80
Precision: 0.00
Recall    : 0.00
F1 Score  : 0.00

```

```
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
```

XGBoost(Bleaching + Longitude)

```
# STEP 1: Imports
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score
from xgboost import XGBClassifier

# STEP 2: Upload dataset
from google.colab import files
uploaded = files.upload()

# STEP 3: Load and preprocess
df = pd.read_csv("dataset.csv")

# Create binary target
df['Collapse'] = df['Species Observed'].apply(lambda x: 1 if x < 100
else 0)

# Encode Bleaching Severity (categorical)
le_bleaching = LabelEncoder()
df['Bleaching Severity'] = le_bleaching.fit_transform(df['Bleaching
Severity'])

# STEP 4: Select features and target
X = df[['Bleaching Severity', 'Longitude']]
y = df['Collapse']

# STEP 5: Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# STEP 6: XGBoost Classifier
xgb = XGBClassifier(use_label_encoder=False, eval_metric='logloss',
random_state=42)
xgb.fit(X_train, y_train)

# STEP 7: Predictions
y_pred = xgb.predict(X_test)

# STEP 8: Evaluation metrics
```

```

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print("\n Evaluation using Bleaching Severity + Longitude:")
print(f"Accuracy : {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall    : {recall:.2f}")
print(f"F1 Score  : {f1:.2f}")

<IPython.core.display.HTML object>

Saving dataset.csv to dataset (25).csv

Evaluation using Bleaching Severity + Longitude:
Accuracy : 0.78
Precision: 0.33
Recall    : 0.10
F1 Score  : 0.15

/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158:
UserWarning: [09:07:07] WARNING: /workspace/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(msg, UserWarning)

```

Random Forest(Bleaching + Longitude)

```

# STEP 0: Imports
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score

# STEP 1: Upload dataset
from google.colab import files
uploaded = files.upload()

# STEP 2: Load and preprocess
df = pd.read_csv("dataset.csv")

# Create binary target
df['Collapse'] = df['Species Observed'].apply(lambda x: 1 if x < 100
else 0)

# Encode Bleaching Severity (categorical)
le_bleaching = LabelEncoder()

```

```

df['Bleaching Severity'] = le_bleaching.fit_transform(df['Bleaching
Severity'])

# STEP 3: Select features and target
X = df[['Bleaching Severity', 'Longitude']]
y = df['Collapse']

# STEP 4: Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# STEP 5: Random Forest Classifier
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)

# STEP 6: Predictions
y_pred = rf.predict(X_test)

# STEP 7: Evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print("\n Evaluation using Bleaching Severity + Longitude:")
print(f"Accuracy : {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall    : {recall:.2f}")
print(f"F1 Score  : {f1:.2f}")

<IPython.core.display.HTML object>

Saving dataset.csv to dataset (26).csv

 Evaluation using Bleaching Severity + Longitude:
Accuracy : 0.75
Precision: 0.31
Recall    : 0.20
F1 Score  : 0.24

```

XGBoost(Location + Latitude)

```

# STEP 1: Imports
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score
from xgboost import XGBClassifier

```



```

# STEP 2: Upload dataset
from google.colab import files
uploaded = files.upload()

# STEP 3: Load and preprocess
df = pd.read_csv("dataset.csv")

# Create binary target
df['Collapse'] = df['Species Observed'].apply(lambda x: 1 if x < 100
else 0)

# Encode Location (categorical)
le_location = LabelEncoder()
df['Location'] = le_location.fit_transform(df['Location'])

# STEP 4: Select features and target
X = df[['Location', 'Latitude']]
y = df['Collapse']

# STEP 5: Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# STEP 6: XGBoost Classifier
xgb = XGBClassifier(use_label_encoder=False, eval_metric='logloss',
random_state=42)
xgb.fit(X_train, y_train)

# STEP 7: Predictions
y_pred = xgb.predict(X_test)

# STEP 8: Evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print("\n Evaluation using Location + Latitude:")
print(f"Accuracy : {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall    : {recall:.2f}")
print(f"F1 Score  : {f1:.2f}")

```

```

Requirement already satisfied: xgboost in
/usr/local/lib/python3.11/dist-packages (2.1.4)
Requirement already satisfied: numpy in
/usr/local/lib/python3.11/dist-packages (from xgboost) (2.0.2)
Requirement already satisfied: nvidia-nccl-cu12 in
/usr/local/lib/python3.11/dist-packages (from xgboost) (2.21.5)

```

Requirement already satisfied: scipy in
/usr/local/lib/python3.11/dist-packages (from xgboost) (1.15.3)

<IPython.core.display.HTML object>

Saving dataset.csv to dataset (27).csv

□ Evaluation using Location + Latitude:

Accuracy : 0.78

Precision: 0.00

Recall : 0.00

F1 Score : 0.00

/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158:

UserWarning: [09:09:27] WARNING: /workspace/src/learner.cc:740:

Parameters: { "use_label_encoder" } are not used.

warnings.warn(smsg, UserWarning)

Random Forest(Location + Latitude)

STEP 0: Imports

import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import LabelEncoder

from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score

STEP 1: Upload dataset

from google.colab import files

uploaded = files.upload()

STEP 2: Load and preprocess

df = pd.read_csv("dataset.csv")

Create binary target

df['Collapse'] = df['Species Observed'].apply(lambda x: 1 if x < 100
else 0)

Encode Location (categorical)

le_location = LabelEncoder()

df['Location'] = le_location.fit_transform(df['Location'])

STEP 3: Select features and target

X = df[['Location', 'Latitude']]

y = df['Collapse']

STEP 4: Train-test split

X_train, X_test, y_train, y_test = train_test_split(X, y,

```

test_size=0.2, random_state=42)

# STEP 5: Random Forest Classifier
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)

# STEP 6: Predictions
y_pred = rf.predict(X_test)

# STEP 7: Evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print("\n Evaluation using Location + Latitude:")
print(f"Accuracy : {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall    : {recall:.2f}")
print(f"F1 Score  : {f1:.2f}")

<IPython.core.display.HTML object>

Saving dataset.csv to dataset (28).csv

Evaluation using Location + Latitude:
Accuracy : 0.76
Precision: 0.17
Recall    : 0.05
F1 Score  : 0.08

```

Comparison chart across all models and group of features

```

# STEP 0: Import libraries
import matplotlib.pyplot as plt
import numpy as np

# STEP 1: Define accuracy values from your experiments
# Replace these with actual accuracy values from your runs
accuracy_all_xgb      = 0.84
accuracy_all_rf       = 0.86
accuracy_sstb_xgb     = 0.84
accuracy_sstb_rf      = 0.87
accuracy_sstl_xgb     = 0.82
accuracy_sstl_rf      = 0.82
accuracy_sstlong_xgb  = 0.84
accuracy_sstlong_rf   = 0.83
accuracy_bl_xgb       = 0.80
accuracy_bl_rf        = 0.80
accuracy_blong_xgb    = 0.78

```

```

accuracy_blong_rf      = 0.75
accuracy_loclong_xgb   = 0.78
accuracy_loclong_rf    = 0.76

# STEP 2: Feature combinations
labels = [
    "SST+Bleaching+Location+Longitude",
    "SST+Bleaching",
    "SST+Location",
    "SST+Longitude",
    "Bleaching+Location",
    "Bleaching+Longitude",
    "Location+Longitude"
]

# XGBoost and Random Forest accuracy values
xgb_scores = [
    accuracy_all_xgb, accuracy_sstb_xgb, accuracy_sstl_xgb,
    accuracy_sstlong_xgb, accuracy_bl_xgb, accuracy_blong_xgb,
    accuracy_loclong_xgb
]

rf_scores = [
    accuracy_all_rf, accuracy_sstb_rf, accuracy_sstl_rf,
    accuracy_sstlong_rf, accuracy_bl_rf, accuracy_blong_rf,
    accuracy_loclong_rf
]

# STEP 3: Plotting
x = np.arange(len(labels))
width = 0.35

plt.figure(figsize=(12, 6))
plt.bar(x - width/2, xgb_scores, width, label='XGBoost',
color='skyblue')
plt.bar(x + width/2, rf_scores, width, label='Random Forest',
color='lightgreen')

# Annotate bars with accuracy values
for i in range(len(labels)):
    plt.text(x[i] - width/2, xgb_scores[i] + 0.005,
f"{xgb_scores[i]:.2f}", ha='center', fontsize=9)
    plt.text(x[i] + width/2, rf_scores[i] + 0.005,
f"{rf_scores[i]:.2f}", ha='center', fontsize=9)

# Plot formatting
plt.ylabel('Accuracy')
plt.xlabel('Feature Combination')
plt.title('Comparison of Accuracy: XGBoost vs Random Forest')
plt.xticks(x, labels, rotation=45, ha='right')

```

```
plt.ylim(0.75, 0.90)
plt.legend()
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```

